

Customer Acquisition & Retention Analysis

AUTHOR

Elizabeth Obst, Shamir Cardenas, Jordan Bona, Olivia
Shipley

PUBLISHED

November 22, 2025

▼ Code

```
library(SMCRM)
```

Warning: package 'SMCRM' was built under R version 4.5.2

▼ Code

```
library(randomForest)
```

Warning: package 'randomForest' was built under R version 4.5.2

randomForest 4.7-1.2

Type rfNews() to see new features/changes/bug fixes.

▼ Code

```
library(rpart)
```

Warning: package 'rpart' was built under R version 4.5.2

▼ Code

```
library(rpart.plot)
```

Warning: package 'rpart.plot' was built under R version 4.5.2

▼ Code

```
library(pdp)
```

Warning: package 'pdp' was built under R version 4.5.2

▼ Code

```
library(pROC)
```

Warning: package 'pROC' was built under R version 4.5.1

Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following objects are masked from 'package:stats':

cov, smooth, var

▼ Code

```
library(corrplot)
```

Warning: package 'corrplot' was built under R version 4.5.1

corrplot 0.95 loaded

▼ Code

```
# Set seed so results are reproducible  
set.seed(123)
```

▼ Code

```
data(acquisitionRetention)  
  
cat("Dataset dimensions:", nrow(acquisitionRetention), "rows,",  
    ncol(acquisitionRetention), "columns\n\n")
```

Dataset dimensions: 500 rows, 15 columns

▼ Code

```
str(acquisitionRetention)
```

```
'data.frame': 500 obs. of 15 variables:  
 $ customer : num 1 2 3 4 5 6 7 8 9 10 ...  
 $ acquisition: num 1 1 1 0 1 1 1 1 0 0 ...  
 $ duration : num 1635 1039 1288 0 1631 ...  
 $ profit   : num 6134 3524 4081 -638 5446 ...  
 $ acq_exp   : num 694 460 249 638 589 ...  
 $ ret_exp   : num 972 450 805 0 920 ...  
 $ acq_exp_sq: num 480998 211628 62016 407644 346897 ...  
 $ ret_exp_sq: num 943929 202077 648089 0 846106 ...  
 $ freq      : num 6 11 21 0 2 7 15 13 0 0 ...  
 $ freq_sq   : num 36 121 441 0 4 49 225 169 0 0 ...  
 $ crossbuy  : num 5 6 6 0 9 4 5 5 0 0 ...  
 $ sow       : num 95 22 90 0 80 48 51 23 0 0 ...  
 $ industry  : num 1 0 0 0 0 1 0 1 0 1 ...  
 $ revenue   : num 47.2 45.1 29.1 40.6 48.7 ...  
 $ employees : num 898 686 1423 181 631 ...
```

EDA

Basic Statistics

▼ Code

```
# Summary statistics
summary(acquisitionRetention)
```

	customer	acquisition	duration	profit
Min.	: 1.0	Min. :0.000	Min. : 0.0	Min. :-1027.0
1st Qu.	:125.8	1st Qu.:0.000	1st Qu.: 0.0	1st Qu.: -316.3
Median	:250.5	Median :1.000	Median : 957.5	Median : 3369.9
Mean	:250.5	Mean :0.676	Mean : 742.5	Mean : 2403.8
3rd Qu.	:375.2	3rd Qu.:1.000	3rd Qu.:1146.2	3rd Qu.: 3931.6
Max.	:500.0	Max. :1.000	Max. :1673.0	Max. : 6134.3
	acq_exp	ret_exp	acq_exp_sq	ret_exp_sq
Min.	: 1.21	Min. : 0.0	Min. :1.460e+00	Min. : 0
1st Qu.	:384.14	1st Qu.: 0.0	1st Qu.:1.476e+05	1st Qu.: 0
Median	:491.66	Median : 398.1	Median :2.417e+05	Median : 158480
Mean	:493.35	Mean : 336.3	Mean :2.712e+05	Mean : 184000
3rd Qu.	:600.21	3rd Qu.: 514.3	3rd Qu.:3.602e+05	3rd Qu.: 264466
Max.	:1027.04	Max. :1095.0	Max. :1.055e+06	Max. :1198937
	freq	freq_sq	crossbuy	sow
Min.	: 0.00	Min. : 0.00	Min. : 0.000	Min. : 0.00
1st Qu.	: 0.00	1st Qu.: 0.00	1st Qu.: 0.000	1st Qu.: 0.00
Median	: 6.00	Median : 36.00	Median : 5.000	Median : 44.00
Mean	: 6.22	Mean : 69.25	Mean : 4.052	Mean : 38.88
3rd Qu.	:11.00	3rd Qu.:121.00	3rd Qu.: 7.000	3rd Qu.: 66.00
Max.	:21.00	Max. :441.00	Max. :11.000	Max. :116.00
	industry	revenue	employees	
Min.	:0.000	Min. :14.49	Min. : 18.0	
1st Qu.	:0.000	1st Qu.:33.53	1st Qu.: 503.0	
Median	:1.000	Median :41.43	Median : 657.5	
Mean	:0.522	Mean :40.54	Mean : 671.5	
3rd Qu.	:1.000	3rd Qu.:47.52	3rd Qu.: 826.0	
Max.	:1.000	Max. :65.10	Max. :1461.0	

▼ Code

```
# Check for missing values
cat("\nMissing values per column:\n")
```

Missing values per column:

▼ Code

```
missing_counts <- colSums(is.na(acquisitionRetention))
print(missing_counts[missing_counts > 0])
```

named numeric(0)

▼ Code

```
if(sum(missing_counts) == 0) {
  cat("No missing values found!\n")
}
```

No missing values found!


Target Variable Analysis

▼ Code

```
par(mfrow = c(1, 2))

# Acquisition counts
acq_table <- table(acquisitionRetention$acquisition)
barplot(acq_table,
        main = "Customer Acquisition Distribution",
        xlab = "Acquired (0=No, 1=Yes)",
        ylab = "Count",
        col = c("coral", "lightgreen"),
        ylim = c(0, max(acq_table) * 1.2))
text(x = c(0.7, 1.9), y = acq_table, labels = acq_table, pos = 3)

# Acquisition percentages
acq_prop <- prop.table(acq_table) * 100
barplot(acq_prop,
        main = "Acquisition Rate (%)",
        xlab = "Acquired (0=No, 1=Yes)",
        ylab = "Percentage",
        col = c("coral", "lightgreen"),
        ylim = c(0, max(acq_prop) * 1.2))
text(x = c(0.7, 1.9), y = acq_prop,
     labels = paste0(round(acq_prop, 1), "%"), pos = 3)
```



▼ Code

```
par(mfrow = c(1, 1))

cat("\nAcquisition Summary:\n")
```

Acquisition Summary:

▼ Code

```
cat("Not Acquired:", acq_table[1], "(", round(acq_prop[1], 1), "%)\n")
```

Not Acquired: 162 (32.4 %)

▼ Code

```
cat("Acquired:", acq_table[2], "(", round(acq_prop[2], 1), "%)\n")
```

Acquired: 338 (67.6 %)

Duration Analysis

▼ Code

```
# Analyze duration for acquired customers
acquired <- acquisitionRetention[acquisitionRetention$acquisition == 1, ]
not_acquired <- acquisitionRetention[acquisitionRetention$acquisition == 0, ]





par(mfrow = c(2, 2))

# Duration for acquired customers
hist(acquired$duration,
  main = "Duration Distribution (Acquired Customers)",
  xlab = "Duration",
  col = "skyblue",
  breaks = 30,
  border = "white")

# Duration with density curve
hist(acquired$duration,
  probability = TRUE,
  main = "Duration with Density Curve",
  xlab = "Duration",
  col = "lightblue",
  breaks = 30,
  border = "white")
lines(density(acquired$duration), col = "red", lwd = 2)

# Boxplot
boxplot(acquired$duration,
  main = "Duration Boxplot",
  ylab = "Duration",
  col = "lightgreen",
  horizontal = FALSE)

# QQ plot to check normality
qqnorm(acquired$duration, main = "Q-Q Plot of Duration")
qqline(acquired$duration, col = "red", lwd = 2)
```

Duration Distribution (Acquired Customer)**Duration with Density Curve****Duration Boxplot****Q-Q Plot of Duration**

▼ Code

```
par(mfrow = c(1, 1))

cat("\nDuration Summary Statistics (Acquired Customers):\n")
```

Duration Summary Statistics (Acquired Customers):

▼ Code

```
print(summary(acquired$duration))
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
654.0	950.5	1072.0	1098.3	1235.8	1673.0

▼ Code

```
cat("\nStandard Deviation:", sd(acquired$duration), "\n")
```

Standard Deviation: 216.713

Feature Distributions


▼ Code

```
# Get all numeric predictors
numeric_cols <- names(acquisitionRetention)[sapply(acquisitionRetention, is.numeric)]
numeric_cols <- setdiff(numeric_cols, c("acquisition", "duration"))

if(length(numeric_cols) > 0) {
  n_plots <- min(12, length(numeric_cols))
  par(mfrow = c(3, 4))

  for(i in 1:n_plots) {
    col_name <- numeric_cols[i]
    hist(acquisitionRetention[[col_name]],
         main = col_name,
         xlab = col_name,
         col = "steelblue",
         breaks = 20,
         border = "white")
  }

  par(mfrow = c(1, 1))
}
```




Features by Acquisition Status

▼ Code

```
# Compare feature distributions between acquired and not acquired
if(length(numeric_cols) > 0) {
  n_plots <- min(9, length(numeric_cols))
  par(mfrow = c(3, 3))

  for(i in 1:n_plots) {
    col_name <- numeric_cols[i]
    boxplot(acquisitionRetention[[col_name]] ~ acquisitionRetention$acquisition,
            main = col_name,
            xlab = "Acquired (0=No, 1=Yes)",
            ylab = col_name,
            col = c("coral", "lightgreen"),
            outline = TRUE)
  }

  par(mfrow = c(1, 1))
}
```




Correlation Analysis

▼ Code

```
# Create correlation matrix
numeric_data <- acquisitionRetention[, c("acquisition", numeric_cols[1:min(15, length(numeric_cols))])]
cor_matrix <- cor(numeric_data, use = "complete.obs")

# Plot correlation heatmap
corrplot(cor_matrix,
         method = "color",
         type = "upper",
         tl.col = "black",
         tl.srt = 45,
         addCoef.col = "black",
         number.cex = 0.7,
         title = "Feature Correlation Matrix",
         mar = c(0, 0, 2, 0))
```

Feature Correlation Matrix



▼ Code

```
# Find highly correlated features with acquisition
cat("\nCorrelation with Acquisition:\n")
```

Correlation with Acquisition:

▼ Code

```
acq_cor <- sort(abs(cor_matrix[, "acquisition"]), decreasing = TRUE)
print(head(acq_cor, 10))
```

	acquisition	profit	ret_exp	crossbuy	sow	freq
acquisition	1.0000000	0.9639299	0.8741157	0.8661546	0.8471508	0.7789100
ret_exp_sq	0.6313023	0.5676057	0.4770195	0.2488373		
employees						
revenue						

Data Preparation

▼ Code

```
# Check the data structure
cat("All columns:\n")
```

All columns:

▼ Code

```
print(names(acquisitionRetention))
```

```
[1] "customer"      "acquisition"    "duration"      "profit"       "acq_exp"
[6] "ret_exp"        "acq_exp_sq"     "ret_exp_sq"    "freq"        "freq_sq"
[11] "crossbuy"       "sow"           "industry"     "revenue"     "employees"
```

▼ Code

```
# Variables to remove (consequences of acquisition, not predictors):
leakage_vars <- c("duration", "profit", "ret_exp", "acq_exp_sq", "ret_exp_sq",
                  "freq", "freq_sq", "crossbuy", "sow")

cat("Removing these variables to prevent leakage:\n")
```

Removing these variables to prevent leakage:

▼ Code

```
print(leakage_vars)
```

```
[1] "duration"      "profit"       "ret_exp"      "acq_exp_sq"   "ret_exp_sq"
[6] "freq"          "freq_sq"     "crossbuy"    "sow"
```

▼ Code

```
# Check for ID columns
id_cols <- grep("id|ID|customer", names(acquisitionRetention),
                 ignore.case = TRUE, value = TRUE)
if(length(id_cols) > 0) {
  cat("\nAlso removing ID columns:", paste(id_cols, collapse = ", "), "\n")
  leakage_vars <- c(leakage_vars, id_cols)
}
```

Also removing ID columns: customer

▼ Code

```
# Create modeling dataset with only legitimate predictors
data_for_models <- acquisitionRetention[, !names(acquisitionRetention) %in% leakage_vars]
```

```
cat("\nVariables Used for Modeling:\n")
```

Variables Used for Modeling:

▼ Code

```
cat("Features:", paste(setdiff(names(data_for_models), "acquisition"), collapse = ", "), "\n")
```

Features: acq_exp, industry, revenue, employees

▼ Code

```
cat("Number of features:", ncol(data_for_models) - 1, "\n")
```

Number of features: 4

▼ Code

```
# Split data 70/30
set.seed(123)
n <- nrow(data_for_models)
train_size <- floor(0.7 * n)
train_indices <- sample(1:n, train_size)

train <- data_for_models[train_indices, ]
test <- data_for_models[-train_indices, ]

cat("\nTraining set:", nrow(train), "observations\n")
```

Training set: 350 observations

▼ Code

```
cat("Test set:", nrow(test), "observations\n")
```

Test set: 150 observations

▼ Code

```
cat("Acquisition rate in training:", round(mean(train$acquisition), 3), "\n")
```

Acquisition rate in training: 0.7

▼ Code

```
cat("Acquisition rate in test:", round(mean(test$acquisition), 3), "\n")
```

Acquisition rate in test: 0.62

Modeling

Random Forest Model

▼ Code

```
# Prepare data (remove duration column)
train_no_duration <- train
test_no_duration <- test

# Build Random Forest
rf_model <- randomForest(
  as.factor(acquisition) ~ .,
  data = train_no_duration,
  ntree = 300,
  importance = TRUE
)

print(rf_model)
```

Call:

```
randomForest(formula = as.factor(acquisition) ~ ., data = train_no_duration,      ntree = 300,
importance = TRUE)
          Type of random forest: classification
                     Number of trees: 300
No. of variables tried at each split: 2
```

OOB estimate of error rate: 21.43%

Confusion matrix:


	0	1	class.error
0	61	44	0.4190476
1	31	214	0.1265306

Variable Importance

▼ Code

```
# Plot variable importance
varImpPlot(rf_model,
           main = "Random Forest Variable Importance",
           n.var = min(20, nrow(importance(rf_model))))
```

Random Forest Variable Importance



▼ Code

```
# Get importance scores
importance_scores <- importance(rf_model)
importance_df <- data.frame(
  Variable = rownames(importance_scores),
  Importance = importance_scores[, "MeanDecreaseGini"]
)
importance_df <- importance_df[order(-importance_df$Importance), ]

cat("\nTop 10 Most Important Variables:\n")
```

Top 10 Most Important Variables:

▼ Code

```
print(head(importance_df, 10))
```

	Variable	Importance
employees	employees	61.512949
acq_exp	acq_exp	41.862147

```
revenue      revenue  36.208863
industry     industry   6.749492
```

Hyperparameter Tuning

▼ Code

```
# Try different mtry values
cat("Tuning mtry parameter\n\n")
```

Tuning mtry parameter

▼ Code

```
results <- data.frame(mtry = numeric(), OOB_error = numeric())

for(m in c(2, 4, 6, 8, 10)) {
  rf_temp <- randomForest(
    as.factor(acquisition) ~.,
    data = train_no_duration,
    ntree = 300,
    mtry = m
  )

  oob_error <- rf_temp$err.rate[300, 1]
  results <- rbind(results, data.frame(mtry = m, OOB_error = oob_error))
  cat("mtry =", m, "| OOB Error =", round(oob_error, 4), "\n")
}
```

mtry = 2 | OOB Error = 0.2171
mtry = 4 | OOB Error = 0.22

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

mtry = 6 | OOB Error = 0.2114

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range


mtry = 8 | OOB Error = 0.2229

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

mtry = 10 | OOB Error = 0.2143

▼ Code

```
# Plot tuning results
plot(results$mtry, results$OOB_error,
  type = "b",
  xlab = "mtry",
  ylab = "OOB Error",
  main = "Random Forest Tuning",
  col = "blue",
  pch = 16,
  lwd = 2)
```



▼ Code

```
best_mtry <- results$mtry[which.min(results$OOB_error)]
cat("\nBest mtry:", best_mtry, "\n")
```

Best mtry: 6

Final Random Forest Model

▼ Code

```
# Build final model with best parameters
rf_final <- randomForest(
  as.factor(acquisition) ~ .,
  data = train_no_duration,
  ntree = 300,
  mtry = best_mtry,
  importance = TRUE
)
```

Warning in randomForest.default(m, y, ...): invalid mtry: reset to within valid range

▼ Code

```
cat("Final Random Forest Model:\n")
```

Final Random Forest Model:

▼ Code

```
print(rf_final)
```

Call:

```
randomForest(formula = as.factor(acquisition) ~ ., data = train_no_duration, ntree = 300,
```

```
mtry = best_mtry, importance = TRUE)
```

```
      Type of random forest: classification
```

```
      Number of trees: 300
```

```
No. of variables tried at each split: 4
```

```
      OOB estimate of error rate: 22.86%
```

Confusion matrix:

	0	1	class.error
0	60	45	0.4285714
1	35	210	0.1428571

	0	1	class.error
0	60	45	0.4285714
1	35	210	0.1428571

▼ Code


```
# Make predictions
rf_pred_class <- predict(rf_final, test_no_duration)
rf_pred_prob <- predict(rf_final, test_no_duration, type = "prob")[, 2]
rf_pred_numeric <- as.numeric(as.character(rf_pred_class))
```

Decision Tree Model

▼ Code

```
# Build decision tree
tree_model <- rpart(
  acquisition ~ .,
  data = train_no_duration,
  method = "class"
)

# Plot complexity parameter
plotcp(tree_model)
```



▼ Code

```
# Prune tree
best_cp <- tree_model$cptable[which.min(tree_model$cptable[, "xerror"]), "CP"]
cat("Best CP:", best_cp, "\n")
```

Best CP: 0.02857143


▼ Code

```
tree_pruned <- prune(tree_model, cp = best_cp)

# Visualize tree
```

```
rpart.plot(tree_pruned,
           main = "Pruned Decision Tree",
           extra = 104,
           fallen.leaves = TRUE)
```

Pruned Decision Tree



▼ Code

```
# Make predictions
tree_pred_prob <- predict(tree_pruned, test_no_duration)[, 2]
tree_pred_class <- ifelse(tree_pred_prob > 0.5, 1, 0)
```

Logistic Regression Model

▼ Code

```
# Build logistic regression
logit_model <- glm(
  acquisition ~.,
  data = train_no_duration,
  family = binomial
)
```

```
summary(logit_model)
```

Call:

```
glm(formula = acquisition ~ ., family = binomial, data = train_no_duration)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-7.6871800	1.0653048	-7.216	5.36e-13 ***
acq_exp	0.0008318	0.0008680	0.958	0.338
industry	1.2727217	0.3073514	4.141	3.46e-05 ***
revenue	0.0869960	0.0161776	5.378	7.55e-08 ***
employees	0.0066901	0.0008758	7.639	2.20e-14 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 427.61 on 349 degrees of freedom

Residual deviance: 285.54 on 345 degrees of freedom

AIC: 295.54

Number of Fisher Scoring iterations: 5

▼ Code

```
# Make predictions
logit_pred_prob <- predict(logit_model, test_no_duration, type = "response")
logit_pred_class <- ifelse(logit_pred_prob > 0.5, 1, 0)
```

Significant Variables

▼ Code

```
# Extract coefficients
coef_summary <- summary(logit_model)$coefficients
significant_vars <- coef_summary[coef_summary[, 4] < 0.05, ]

cat("Significant Variables (p < 0.05):\n")
```

Significant Variables (p < 0.05):

▼ Code

```
print(significant_vars)
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-7.687180040	1.0653048069	-7.215944	5.356102e-13

industry	1.272721705	0.3073513643	4.140934	3.458945e-05
revenue	0.086995954	0.0161775551	5.377571	7.549738e-08
employees	0.006690118	0.0008758244	7.638652	2.195082e-14


▼ Code

```
# Plot coefficients
if(nrow(significant_vars) > 1) {
  sig_coefs <- significant_vars[-1, 1] # Remove intercept

  # Sort by absolute value
  sig_coefs_sorted <- sort(sig_coefs)

  par(mar = c(5, 8, 4, 2))
  barplot(sig_coefs_sorted,
          horiz = TRUE,
          las = 1,
          col = ifelse(sig_coefs_sorted > 0, "steelblue", "coral"),
          main = "Significant Logistic Regression Coefficients",
          xlab = "Coefficient Value")
  abline(v = 0, lty = 2)
  par(mar = c(5, 4, 4, 2))
}
```

Significant Logistic Regression Coefficients



Model Comparison

▼ Code

```
# Calculate accuracy
calc_accuracy <- function(actual, predicted) {
  sum(actual == predicted) / length(actual)
}

# Sanity check for data leakage
cat("Diagnostic Check:\n")
```

Diagnostic Check:

▼ Code

```
cat("Test set size:", nrow(test), "\n")
```

Test set size: 150

▼ Code

```
cat("Actual acquisition distribution:", table(test$acquisition), "\n")
```

Actual acquisition distribution: 57 93

▼ Code

```
cat("RF predictions distribution:", table(rf_pred_numeric), "\n\n")
```

RF predictions distribution: 40 110

▼ Code

```
# Calculate metrics
rf_accuracy <- calc_accuracy(test$acquisition, rf_pred_numeric)
tree_accuracy <- calc_accuracy(test$acquisition, tree_pred_class)
logit_accuracy <- calc_accuracy(test$acquisition, logit_pred_class)

cat("Accuracy Results:\n")
```

Accuracy Results:

▼ Code

```
cat("Random Forest:", round(rf_accuracy, 4), "\n")
```

Random Forest: 0.7667

▼ Code

```
cat("Decision Tree:", round(tree_accuracy, 4), "\n")
```

Decision Tree: 0.7333

▼ Code

```
cat("Logistic Regression:", round(logit_accuracy, 4), "\n\n")
```

Logistic Regression: 0.8

▼ Code

```
if(max(rf_accuracy, tree_accuracy, logit_accuracy) > 0.95) {  
  cat("Note: Very high accuracy detected. This could indicate:\n")  
  cat("- Strong predictive patterns in the data\n")  
  cat("- Potential overfitting (validate on new data)\n\n")  
}
```

Confusion Matrices

▼ Code

```
par(mfrow = c(1, 3))  
  
# Function to create nice confusion matrix  
plot_confusion <- function(actual, predicted, title) {  
  cm <- table(Predicted = predicted, Actual = actual)  
  
  # Calculate percentages  
  cm_pct <- prop.table(cm, 2) * 100  
  
  # Create color matrix  
  colors <- matrix(c("lightcoral", "lightgreen", "lightgreen", "lightcoral"),  
    nrow = 2)  
  
  # Plot  
  plot(0, 0, type = "n", xlim = c(0, 2), ylim = c(0, 2),  
    xlab = "Actual", ylab = "Predicted", main = title,  
    xaxt = "n", yaxt = "n")  
  
  # Add colored rectangles  
  rect(0, 0, 1, 1, col = colors[1, 1], border = "black")  
  rect(1, 0, 2, 1, col = colors[1, 2], border = "black")
```

```


rect(0, 1, 1, 2, col = colors[2, 1], border = "black")
rect(1, 1, 2, 2, col = colors[2, 2], border = "black")

# Add text
text(0.5, 1.5, paste0(cm[2, 1], "\n(", round(cm_pct[2, 1], 1), "%)"), cex = 1.2)
text(1.5, 1.5, paste0(cm[2, 2], "\n(", round(cm_pct[2, 2], 1), "%)"), cex = 1.2)
text(0.5, 0.5, paste0(cm[1, 1], "\n(", round(cm_pct[1, 1], 1), "%)"), cex = 1.2)
text(1.5, 0.5, paste0(cm[1, 2], "\n(", round(cm_pct[1, 2], 1), "%)"), cex = 1.2)

# Add axis labels
axis(1, at = c(0.5, 1.5), labels = c("0", "1"))
axis(2, at = c(0.5, 1.5), labels = c("0", "1"))
}

plot_confusion(test$acquisition, rf_pred_numeric, "Random Forest")
plot_confusion(test$acquisition, tree_pred_class, "Decision Tree")
plot_confusion(test$acquisition, logit_pred_class, "Logistic Regression")

```



▼ Code

```
par(mfrow = c(1, 1))
```


ROC Curves

▼ Code

```
# Calculate ROC curves
rf_roc <- roc(test$acquisition, rf_pred_prob, quiet = TRUE)
tree_roc <- roc(test$acquisition, tree_pred_prob, quiet = TRUE)
logit_roc <- roc(test$acquisition, logit_pred_prob, quiet = TRUE)

# Plot
plot(rf_roc, col = "blue", lwd = 2, main = "ROC Curves Comparison")
plot(tree_roc, col = "red", lwd = 2, add = TRUE)
plot(logit_roc, col = "green", lwd = 2, add = TRUE)
abline(a = 0, b = 1, lty = 2, col = "gray")

legend("bottomright",
       legend = c(
         paste("Random Forest (AUC =", round(auc(rf_roc), 3), ")"),
         paste("Decision Tree (AUC =", round(auc(tree_roc), 3), ")"),
         paste("Logistic Regression (AUC =", round(auc(logit_roc), 3), ")")
       ),
       col = c("blue", "red", "green"),
       lwd = 2,
       cex = 0.9)
```



Performance Summary Table

▼ Code

```
# Create summary table
summary_df <- data.frame(
  Model = c("Random Forest", "Decision Tree", "Logistic Regression"),
  Accuracy = c(rf_accuracy, tree_accuracy, logit_accuracy),
  AUC = c(auc(rf_roc), auc(tree_roc), auc(logit_roc))
)
print(summary_df)
```

	Model	Accuracy	AUC
1	Random Forest	0.7666667	0.8206942
2	Decision Tree	0.7333333	0.7425957
3	Logistic Regression	0.8000000	0.8490851

▼ Code

```
best_model <- summary_df$Model[which.max(summary_df$AUC)]
```

```
cat("\nBest Model:", best_model, "\n")
```

Best Model: Logistic Regression

Duration Prediction

▼ Code

```
# For duration prediction, we use the full dataset (including retention variables)
# because duration is only predicted after acquisition

# Get acquired customers from original data
train_acquired_full <- acquisitionRetention[train_indices, ]
train_acquired_full <- train_acquired_full[train_acquired_full$acquisition == 1, ]

test_acquired_full <- acquisitionRetention[-train_indices, ]
test_acquired_full <- test_acquired_full[test_acquired_full$acquisition == 1, ]

cat("Acquired customers in training:", nrow(train_acquired_full), "\n")
```

Acquired customers in training: 245

▼ Code

```
cat("Acquired customers in test:", nrow(test_acquired_full), "\n\n")
```

Acquired customers in test: 93

▼ Code

```
if(nrow(train_acquired_full) > 10) {
  # Remove acquisition column and any ID columns for duration model
  duration_predictors <- setdiff(names(train_acquired_full),
                                    c("acquisition", "duration", id_cols))

  cat("Using these predictors for duration:\n")
  print(duration_predictors)
  cat("\n")

  # Build duration model with all available features
  rf_duration <- randomForest(
    duration ~ .,
    data = train_acquired_full[, c("duration", duration_predictors)],
    ntree = 300,
    importance = TRUE
  )
```

```

print(rf_duration)

# Variable importance for duration
varImpPlot(rf_duration,
            main = "Important Variables for Duration Prediction",
            n.var = min(15, nrow(importance(rf_duration)))) 

# Predictions
if(nrow(test_acquired_full) > 0) {
  duration_pred <- predict(rf_duration, test_acquired_full)

  # Performance metrics
  rmse <- sqrt(mean((test_acquired_full$duration - duration_pred)^2))
  mae <- mean(abs(test_acquired_full$duration - duration_pred))
  r_squared <- cor(test_acquired_full$duration, duration_pred)^2

  cat("\nDuration Model Performance:\n")
  cat("RMSE:", round(rmse, 3), "\n")
  cat("MAE:", round(mae, 3), "\n")
  cat("R-squared:", round(r_squared, 3), "\n")

  # Plot actual vs predicted
  par(mfrow = c(1, 2))

  plot(test_acquired_full$duration, duration_pred,
        xlab = "Actual Duration",
        ylab = "Predicted Duration",
        main = "Actual vs Predicted Duration",
        pch = 16,
        col = rgb(0, 0, 1, 0.5))
  abline(0, 1, col = "red", lwd = 2)

  # Residuals plot
  residuals <- test_acquired_full$duration - duration_pred
  plot(duration_pred, residuals,
        xlab = "Predicted Duration",
        ylab = "Residuals",
        main = "Residual Plot",
        pch = 16,
        col = rgb(1, 0, 0, 0.5))
  abline(h = 0, col = "blue", lwd = 2, lty = 2)

  par(mfrow = c(1, 1))
}
}
}

```

Using these predictors for duration:

```

[1] "profit"      "acq_exp"     "ret_exp"     "acq_exp_sq"  "ret_exp_sq"
[6] "freq"        "freq_sq"     "crossbuy"    "sow"         "industry"
[11] "revenue"     "employees"

```

Call:

```
randomForest(formula = duration ~ ., data = train_acquired_full[, c("duration",
duration_predictors)], ntree = 300, importance = TRUE)
```

Type of random forest: regression


Number of trees: 300

No. of variables tried at each split: 4

Mean of squared residuals: 1489.109

% Var explained: 96.83

Important Variables for Duration Prediction




Duration Model Performance:

RMSE: 32.878

MAE: 24.195

R-squared: 0.978



▼ Code

```
# Get importance scores
importance_scores <- importance(rf_duration)
print(importance_scores)
```

	%IncMSE	IncNodePurity
profit	11.9160609	2094958.36
acq_exp	3.5378055	109569.96
ret_exp	19.9865764	4121608.24
acq_exp_sq	2.9413500	111315.55
ret_exp_sq	21.5385244	4215158.77
freq	13.1692561	149985.52
freq_sq	12.7782403	180242.94
crossbuy	0.4913718	72877.19
sow	0.8812465	71679.23
industry	1.0095582	13874.29
revenue	1.3968278	116799.95
employees	2.6180476	155974.15

▼ Code

```
# Sort by %IncMSE
```

```
importance_scores[order(importance_scores[,1], decreasing = TRUE), ]
```

	%IncMSE	IncNodePurity
ret_exp_sq	21.5385244	4215158.77
ret_exp	19.9865764	4121608.24
freq	13.1692561	149985.52
freq_sq	12.7782403	180242.94
profit	11.9160609	2094958.36
acq_exp	3.5378055	109569.96
acq_exp_sq	2.9413500	111315.55
employees	2.6180476	155974.15
revenue	1.3968278	116799.95
industry	1.0095582	13874.29
sow	0.8812465	71679.23
crossbuy	0.4913718	72877.19

Partial Dependence Plots

▼ Code

```
# Get top important variables
imp <- importance(rf_final)
top_vars <- head(rownames(imp)[order(-imp[, 1])], 6)

# Store PDP data for interpretation
pdp_summary <- list()

par(mfrow = c(2, 3))
for(var in top_vars) {
  pdp_data <- partial(rf_final,
                      pred.var = var,
                      prob = TRUE,
                      which.class = "1")


  # Store summary statistics
  pdp_summary[[var]] <- data.frame(
    variable = var,
    min_value = min(pdp_data[[1]]),
    max_value = max(pdp_data[[1]]),
    min_prob = min(pdp_data$yhat),
    max_prob = max(pdp_data$yhat),
    range_effect = max(pdp_data$yhat) - min(pdp_data$yhat),
    mean_prob = mean(pdp_data$yhat)
  )

  # Plot
  plot(pdp_data[[1]], pdp_data$yhat,
        type = "l",
        lwd = 2,
```

```

    col = "blue",
    main = paste("Effect of", var),
    xlab = var,
    ylab = "Probability of Acquisition")
grid()
}
par(mfrow = c(1, 1))

```



▼ Code

```

# Print summary for interpretation
pdp_df <- do.call(rbind, pdp_summary)
print(pdp_df)

```

	variable	min_value	max_value	min_prob	max_prob	range_effect
employees	employees	18.00	1461.00	0.2121143	0.8997714	0.6876571
acq_exp	acq_exp	1.21	970.31	0.3454095	0.8100476	0.4646381
revenue	revenue	16.29	65.10	0.5105524	0.8494381	0.3388857
industry	industry	0.00	1.00	0.6350000	0.7494381	0.1144381
	mean_prob					
employees	0.6747970					
acq_exp	0.5580472					
revenue	0.6961122					
industry	0.6922190					

Key Findings

Model Performance

▼ Code

```
cat("Final Model Comparison:\n\n")
```

Final Model Comparison:

▼ Code

```
print(summary_df)
```

	Model	Accuracy	AUC
1	Random Forest	0.7666667	0.8206942
2	Decision Tree	0.7333333	0.7425957
3	Logistic Regression	0.8000000	0.8490851

▼ Code

```
cat("\nBest performing model:", best_model, "\n")
```

Best performing model: Logistic Regression

▼ Code

```
cat("Best AUC:", round(max(summary_df$AUC), 4), "\n")
```

Best AUC: 0.8491

▼ Code

```
cat("Best Accuracy:", round(max(summary_df$Accuracy), 4), "\n")
```

Best Accuracy: 0.8

Important Insights

1. The best model was Logistic Regression - Best AUC: 0.8491 Best Accuracy: 0.8
2. The top most important variables for predicting acquisition are: employees, acq_exp, revenue, and industry

▼ Code

```
print(head(importance_df, 5))
```

```
Variable Importance
employees employees 61.512949
acq_exp acq_exp 41.862147
revenue revenue 36.208863
industry industry 6.749492
```

3. Duration model performance: Random Forest RMSE: 32.878 MAE: 24.195 R-squared: 0.978